

Clustering

Clustering

- Process of grouping a set of examples (samples)
- Clustering generates a **partition** consisting of **cohesive groups or clusters** from given collection of examples (samples)



- **For example:**
 - Grouping students in a class based on gender
 - Grouping students in a class based the month of birth
 - Grouping the students based on the place of sitting

Clustering

- Process of grouping a set of examples
- Clustering generates a **partition** consisting of **cohesive groups or clusters** from given collection of examples



- The examples to be clustered are either **labelled** or **unlabelled**
 - **Algorithms which cluster labelled examples:**
 - Supervised clustering
 - Classification: **Learning by examples**
 - **Algorithms which cluster unlabelled examples:**
 - Unsupervised clustering
 - Do not rely on predefined classes
 - **Learning by observation**, rather than learning by examples.

3

Clustering

- Clustering is a two step process
 - Step1: **Partition the collection of examples (clustering)**
 - Learning by observation (training phase)
 - **Group the collection of examples into finite number of clusters** such that the examples that are **similar** to one another within the same cluster and are **dissimilar** to examples in other clusters
 - Obtaining cluster labels
 - **Unsupervised learning:** Do not rely on predefined classes and class-labelled training examples
 - Step2: **Assign cluster labels to examples**
 - Testing phase

4

Categorization of Clustering Methods

- Partitioning methods
- Hierarchical methods
- Density-based methods

5

Categorization of Clustering Methods

- Partitioning methods:
 - These methods construct K partitions of the data, where each partition represents a cluster
 - **Idea**: Cluster the collection of examples based on the **distance between examples**
 - Results in spherical shaped cluster
 - 1. K -means algorithm
 - 2. K -medoids algorithm
 - 3. Gaussian mixture model
- Hierarchical methods:
 - These methods create a **hierarchical decomposition** of the collection of examples
 - Results in spherical shaped cluster
 - 1. Agglomerative approach (bottom-up approach)
 - 2. Divisive approach (top-down approach)

6

Categorization of Clustering Methods

- Density-based methods:
 - These methods cluster collection of examples based on the notion of density
 - General idea: To continue growing the given cluster as long as density (number of examples) in the neighbourhood exceeds some threshold
 - Example:
 - DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

7

Partitioning Method based Clustering

Classical Portioning Methods

- **Centroid-based technique:**
 - Partition the collection of examples into K clusters based on the distance between examples
 - Cluster similarity is measured in regard to the **sample mean** of the examples within a cluster
 - **Cluster centroid** or **center of gravity**: Sample mean value of the examples within a cluster
 - Cluster center is used to represent the cluster
 - Example: K -means algorithm
- **Representative object-based technique:**
 - **Actual example** is considered to represent the cluster
 - One representative example per cluster
 - Example: K -medoids algorithm

9

K -Means Clustering Algorithm

- Dividing the data into K groups or partitions
- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$, $\mathbf{x}_n \in \mathbb{R}^d$ and K
- **Target:** Partition the set \mathcal{D} into K clusters (disjoint subsets), $\{\mathcal{D}_k\}_{k=1}^K$
 - Each of the clusters is associated with centers, μ_k , $k=1, 2, \dots, K$
 - **Come up with the centers of clusters**
 - Cluster center acts as a **cluster representative**
- Euclidean distance with center of a cluster can be used as a measure of dissimilarity

10

K-Means Clustering Algorithm: Training Phase

- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$, $\mathbf{x}_n \in \mathbb{R}^d$ and K
- 1. Initialize the cluster center, $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$ using randomly selected K data points in \mathcal{D}
- 2. Assign each data point \mathbf{x}_n to cluster center k^*

$$k^* = \arg \min_k \|\mathbf{x}_n - \boldsymbol{\mu}_k\|^2 \quad \text{Squared Euclidian distance}$$
- 3. **Update** $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$: Re-compute $\boldsymbol{\mu}_k$ after assigning all the data points.

$$\hat{\boldsymbol{\mu}}_k = \frac{\sum_{\mathcal{D}_k} \mathbf{x}_n}{N_k}$$

\mathcal{D}_k : Data for cluster k

N_k : Number of examples in cluster k
- 4. Repeat the steps 2 and 3 until the convergence

11

K-Means Clustering Algorithm: Training Phase


- **Convergence criteria:**
 - No change in the cluster assignment **OR**
 - The difference between the **distortion measure (J)** in the successive iteration falls below the threshold
 - **Distortion measure (J)** : Sum of the squares of the distance of each example to its assigned cluster center

$$J = \sum_{n=1}^N \sum_{k=1}^K z_{nk} \|\mathbf{x}_n - \boldsymbol{\mu}_k\|^2$$

z_{nk} is 1 if \mathbf{x}_n belongs to cluster k , otherwise 0

12

Illustration of K -Means Clustering



- Boundary between the cluster is linear
- **Hard clustering**: Each example must belong to exactly one group

13

Modified K -Means Clustering Algorithm

- Dividing the data into K groups or partitions
- **Given**: Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^d$ and K
- **Target**: Partition the set \mathcal{D} into K clusters (disjoint subsets), $\{\mathcal{D}_k\}_{k=1}^K$
 - Each of the clusters is associated with centers, $\boldsymbol{\mu}_k, k=1, 2, \dots, K$
 - **Better representative for a cluster**
 - Come up with the **centers of clusters** and **variance & covariance (covariance matrix)**
 - Cluster center and covariance matrix act as **cluster representatives**
 - For $d=2$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \boldsymbol{\mu} = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} = \begin{bmatrix} E[x_1] \\ E[x_2] \end{bmatrix}$$

$$\boldsymbol{\Sigma} = \begin{bmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{21} & \sigma_2^2 \end{bmatrix} \quad \boldsymbol{\Sigma} = \begin{bmatrix} E[(x_1 - \mu_1)^2] & E[(x_1 - \mu_1)(x_2 - \mu_2)] \\ E[(x_2 - \mu_2)(x_1 - \mu_1)] & E[(x_2 - \mu_2)^2] \end{bmatrix}$$
- **Mahalanobis distance** with cluster representatives can be used as a measure of dissimilarity

14

Modified K -Means Clustering Algorithm: Training Phase

- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$, $\mathbf{x}_n \in \mathbb{R}^d$ and K

1. Initialize the cluster center, $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$ using randomly selected K data points in \mathcal{D}
2. Initialize the covariance matrix, $\boldsymbol{\Sigma}_k$, $k=1, 2, \dots, K$ using unit matrix
3. Assign each data point \mathbf{x}_n to cluster center k^*

$$k^* = \arg \min_k (\mathbf{x}_n - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k) \quad \text{Squared Mahalanobis distance}$$

4. **Update $\boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$** , $k=1, 2, \dots, K$: Re-compute $\boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$ after assigning all the data points.

$$\hat{\boldsymbol{\mu}}_k = \frac{\sum_{\mathcal{D}_k} \mathbf{x}_n}{N_k} \quad \hat{\boldsymbol{\Sigma}}_k = \frac{\sum_{\mathcal{D}_k} (\mathbf{x}_n - \hat{\boldsymbol{\mu}}_k)(\mathbf{x}_n - \hat{\boldsymbol{\mu}}_k)^\top}{N_k} \quad \begin{array}{l} \mathcal{D}_k: \text{Data for cluster } k \\ N_k: \text{Number of} \\ \text{examples in cluster } k \end{array}$$

5. Repeat the steps 3 and 4 until the convergence

15

Modified K -Means Clustering Algorithm: Training Phase

- **Convergence criteria:**
 - No change in the cluster assignment **OR**
 - The difference between the **distortion measure (J)** in the successive iteration falls below the threshold


$$J = \sum_{n=1}^N \sum_{k=1}^K z_{nk} \left[(\mathbf{x}_n - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k) \right]$$

z_{nk} is 1 if \mathbf{x}_n belongs to cluster k , otherwise 0

- **Hard clustering:** Each example must belong to exactly one group

16

Illustration of K -Means Clustering




- Boundary between the cluster is quadratic
- **Hard clustering**: Each example must belong to exactly one group

17


Illustration of K -Means Clustering

$K = 3$



Measure of
Dissimilarity:

**Euclidian
Distance**




**Mahalanobis
Distance**

18

Elbow Method to Choose K

- Determine the **distortion measure** for different values of K
- Plot the **K vs Distortion**
- **Optimal number of clusters**: Select the value of K at the “elbow” i.e. the point after which the distortion start decreasing in a linear fashion



19

K -Means Clustering Algorithm: Training Phase

- **Given**: Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$, $\mathbf{x}_n \in \mathbb{R}^d$ and K
1. Initialize the cluster center, $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$ using randomly selected K data points in \mathcal{D}
 2. Assign each data point \mathbf{x}_n to cluster center k^*

$$k^* = \arg \min_k \|\mathbf{x}_n - \boldsymbol{\mu}_k\|^2 \quad \text{Squared Euclidian distance}$$
 3. **Update** $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$: Re-compute $\boldsymbol{\mu}_k$ after assigning all the data points.


$$\hat{\boldsymbol{\mu}}_k = \frac{\sum_{\mathbf{x}_n \in \mathcal{D}_k} \mathbf{x}_n}{N_k}$$

\mathcal{D}_k : Data for cluster k

N_k : Number of examples in cluster k
 4. Repeat the steps 2 and 3 until the convergence

20


Illustration of K -Means Clustering



- Boundary between the cluster is linear
- **Hard clustering**: Each example must belong to exactly one group

21

Soft Clustering



- **Soft clustering**: Each example belong to each group with some probability
 - Fuzzyness at the boundary of the clusters
- Gaussian mixture model (GMM) is one of the soft clustering techniques
- GMM can be seen as similar to K -means clustering
- Each cluster is represented as Gaussian density

22

Gaussian Mixture Model (GMM)

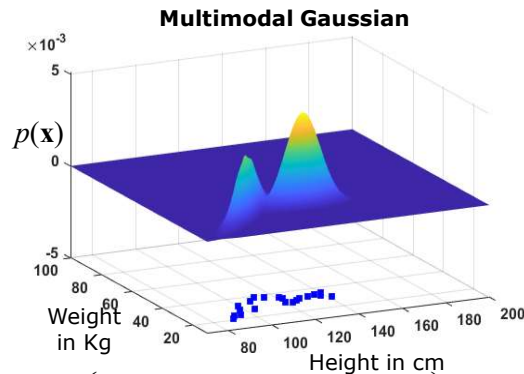
- Data is considered to have **multiple clusters** and each cluster is an Gaussian distribution
- Given:** Training data having N samples

$$\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n, \dots, \mathbf{x}_N\}, \quad \mathbf{x}_n \in \mathbb{R}^d$$

- GMM is a **linear superposition** of multiple **Gaussian components**:

$$p(\mathbf{x}) = \sum_{k=1}^K w_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

$$\mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}_k|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1}(\mathbf{x} - \boldsymbol{\mu}_k)\right)$$



23

Gaussian Mixture Model (GMM)

- GMM is a **linear superposition** of multiple Gaussians:

$$p(\mathbf{x}) = \sum_{k=1}^K w_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- For a d -dimensional feature vector representation of data, the **parameters** of GMM are
 - Mixture coefficients**, w_k , $k = 1, 2, \dots, K$
 - Mixture weight or Strength of each clusters (or mixtures or modes)
 - Property: $\sum_{k=1}^K w_k = 1$
 - d -dimensional **mean vector**, $\boldsymbol{\mu}_k$, $k = 1, 2, \dots, K$
 - $d \times d$ size **covariance matrices**, $\boldsymbol{\Sigma}_k$, $k = 1, 2, \dots, K$
- Training process objective:**
 - Partition the data into K groups
 - To estimate the parameters of the each cluster in GMM

24

Expectation-Maximization (EM) for GMMs

- Given a Gaussian mixture model, the goal is to maximize the likelihood function with respect to the parameters
- Given: Training data having N samples

$$\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n, \dots, \mathbf{x}_N\}, \quad \mathbf{x}_n \in \mathbb{R}^d$$

- Initialize the mean vectors $\boldsymbol{\mu}_k$, covariance matrices $\boldsymbol{\Sigma}_k$ and mixing coefficients w_k , and evaluate the initial value of the log likelihood
 - Initialize the cluster center, $\boldsymbol{\mu}_k$, $k=1, 2, \dots, K$ using randomly selected K data points in \mathcal{D}
 - Initialize the covariance matrix, $\boldsymbol{\Sigma}_k$, $k=1, 2, \dots, K$ using unit matrix
 - Initialize the mixing coefficient $w_k = \frac{1}{K}$, $k=1, 2, \dots, K$

25

Expectation-Maximization (EM) for GMMs

- Given a Gaussian mixture model, the goal is to maximize the likelihood function with respect to the parameters
- Given: Training data having N samples

$$\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n, \dots, \mathbf{x}_N\}, \quad \mathbf{x}_n \in \mathbb{R}^d$$

- Initialize the mean vectors $\boldsymbol{\mu}_k$, covariance matrices $\boldsymbol{\Sigma}_k$ and mixing coefficients w_k , and evaluate the initial value of the log likelihood

$$p(\mathcal{D} | \boldsymbol{\theta}) = \prod_{n=1}^N p(\mathbf{x}_n | \boldsymbol{\theta}) \quad \text{where } \boldsymbol{\theta} = [w_1 \dots w_K, \boldsymbol{\mu}_1 \dots \boldsymbol{\mu}_K, \boldsymbol{\Sigma}_1 \dots \boldsymbol{\Sigma}_K]^\top$$

$$\mathcal{L}(\boldsymbol{\theta}) = \ln p(\mathcal{D} | \boldsymbol{\theta}) = \sum_{n=1}^N \ln p(\mathbf{x}_n | \boldsymbol{\theta}) = \sum_{n=1}^N \ln \left(\sum_{k=1}^K w_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right)$$

- E-step:** Evaluate the responsibilities $\gamma_k(\mathbf{x})$ using the current parameter values – Assign the data points to each cluster

26


EM Method – Responsibility Term

- A quantity that plays an important role is the **responsibility term**, $\gamma_k(\mathbf{x})$

- It is given by

$$\gamma_k(\mathbf{x}) = \frac{w_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{k=1}^K w_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

- w_k : **mixture coefficient** or **prior probability of cluster k** ,
- $\gamma_k(\mathbf{x})$ gives the **posterior probability of the cluster k** for the observation \mathbf{x}



27

Expectation-Maximization (EM) for GMMs

- Given a Gaussian mixture model, the goal is to **maximize the likelihood function with respect to the parameters**

1. Initialize the **mean vectors $\boldsymbol{\mu}_k$** , **covariance matrices $\boldsymbol{\Sigma}_k$** and **mixing coefficients w_k** , and evaluate the initial value of the log likelihood

2. **E-step**: Evaluate the responsibilities $\gamma_k(\mathbf{x})$ using the current parameter values

3. **M-step**: Re-estimate the parameters $\boldsymbol{\mu}_k^{new}$, $\boldsymbol{\Sigma}_k^{new}$ and w_k^{new} using the current responsibilities

$$\boldsymbol{\mu}_k^{new} = \frac{1}{N_k} \sum_{n=1}^N \gamma_k(\mathbf{x}_n) \mathbf{x}_n$$

$$\boldsymbol{\Sigma}_k^{new} = \frac{1}{N_k} \sum_{n=1}^N \gamma_k(\mathbf{x}_n) (\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^T$$

$$w_k^{new} = \frac{N_k}{N}$$

$$N_k = \sum_{n=1}^N \gamma_k(\mathbf{x}_n)$$

- N_k : Effective number of points assigned to the cluster k

28


Expectation-Maximization (EM) for GMMs

- Given a Gaussian mixture model, the goal is to maximize the likelihood function with respect to the parameters
 1. Initialize the mean vectors μ_k , covariance matrices Σ_k and mixing coefficients w_k , and evaluate the initial value of the log likelihood
 2. **E-step**: Evaluate the responsibilities $\gamma_k(\mathbf{x})$ using the current parameter values
 3. **M-step**: Re-estimate the parameters μ_k^{new} , Σ_k^{new} and w_k^{new} using the current responsibilities
 4. Evaluate the log likelihood and check for convergence of the log likelihood
 - If the convergence criterion is not satisfied return to step 2

29

Expectation-Maximization (EM) for GMMs

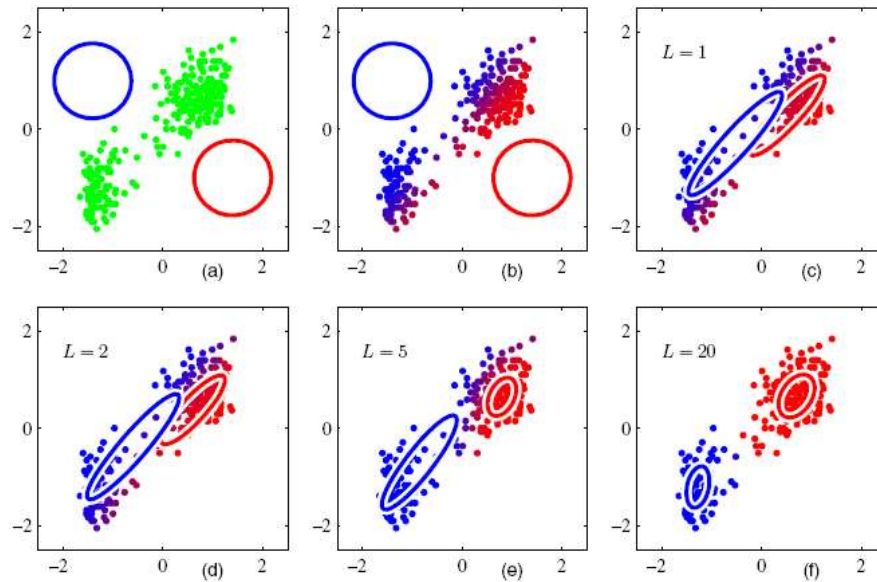
- Convergence criterion**: Difference between log likelihoods of successive iterations fall below a threshold (E.g. 10^{-3})



$$\mathcal{L}(\theta) = \ln p(\mathcal{D} | \theta) = \sum_{n=1}^N \ln \left(\sum_{k=1}^K w_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k) \right)$$

30

Illustration of Parameter Estimation




C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006.

31

Elbow Method to Choose K

- Determine the **total data log likelihood** for different values of K
- Plot the K vs **total data log likelihood**
- **Optimal number of clusters**: Select the value of K at the "elbow" i.e. the point after which the log likelihood start **increasing in a linear fashion**



32


Application of Clustering: Segmentation using Clustering



33

Application of Clustering: Cell and Nucleus Segmentation

- An image is divided into patches of size 7 x 7
- From each patch, **mean** and **variance** of pixel values is considered as features
- 100 images with different numbers of cells are considered for training




K-Means

GMM

34

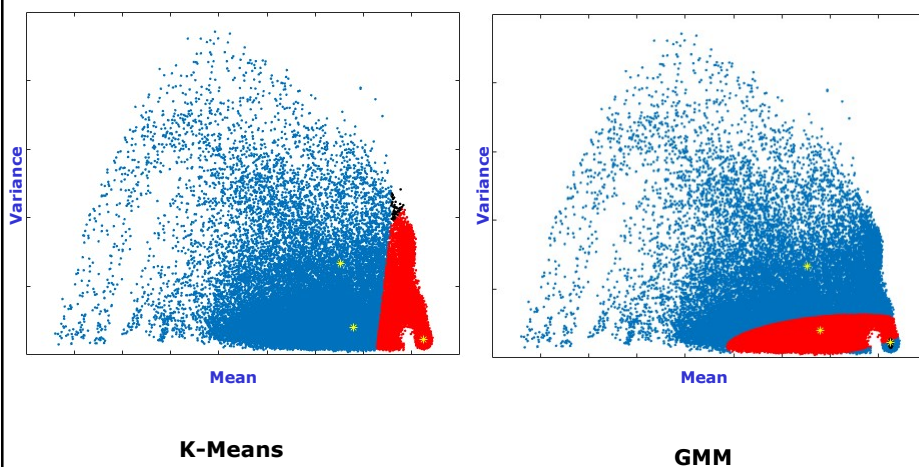
Application of Clustering: Cell and Nucleus Segmentation

- An image is divided into patches of size 7×7
- From each patch, **mean** and **variance** of pixel values is considered as features
- 100 images with different numbers of cells are considered for training



35

Hard Clustering vs Soft Clustering



36

K-Medoid Clustering Algorithms

- Related to K -means clustering
- The K -means algorithm is sensitive to outliers because an example with extremely large value may substantially distort the distribution of data
- Solution: One of the data points is chosen as representative of cluster, instead of mean value of the cluster
- It replaces the means of cluster with modes
- Partitioning around medoids
- A medoid of a finite dataset: The data point from the set, whose average dissimilarity (distance) to all the points is minimal
 - The most centrally located point in the set

37

K-Medoid Clustering Algorithm

- Given: Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$, $\mathbf{x}_n \in \mathbb{R}^d$ and K
1. Initialize the medoid, $\hat{\mathbf{x}}_k$, $k=1, 2, \dots, K$ using randomly selected K data points in \mathcal{D}
 2. Assign each data point \mathbf{x}_n to the closest medoid

$$k^* = \arg \min_k \|\mathbf{x}_n - \hat{\mathbf{x}}_k\|^2 \quad \text{Squared Euclidian distance}$$
 3. Update medoids $\hat{\mathbf{x}}_k$, $k=1, 2, \dots, K$
 - For each data point \mathbf{x}_n assigned to a cluster k compute the average dissimilarity (distance) of \mathbf{x}_n to all the data points assigned to cluster k

$$\text{Average dissimilarity for } \mathbf{x}_n = \frac{\sum_{\mathbf{x}_m \in \mathcal{D}_k} \|\mathbf{x}_n - \mathbf{x}_m\|^2}{N_k} \quad \begin{array}{l} N_k: \text{Number of} \\ \text{examples in cluster } k \end{array}$$
 - Select the example with minimum average dissimilarity as medoid
 4. Repeat the steps 2 and 3 until the convergence

38

K-Medoid Clustering Algorithm

- Convergence criteria:
 - No change in the cluster assignment **OR**
 - The difference between the distortion measure (absolute-error) (J) in the successive iteration falls below the threshold
 - Distortion measure (J) : Sum of the squares of the distance of each example to its corresponding reference point (medoid)

$$J = \sum_{n=1}^N \sum_{k=1}^K z_{nk} \|\mathbf{x}_n - \hat{\mathbf{x}}_k\|^2$$

z_{nk} is 1 if \mathbf{x}_n belongs to cluster k , otherwise 0

- Optimal number of clusters (k) can be obtained using elbow method

39

Evaluation of Clustering: Purity Score

- Let us assume that class index for each example is given
- Purity score: Purity is a measure of the extent to which clusters contain a single class
 - For each cluster, count the number of data points from the most common class
 - Take the sum over all clusters and divide by the total number of data points
- Let M be the number of classes, $C_1, C_2, \dots, C_m, \dots, C_M$
- Let K be the number of clusters, $k = 1, 2, \dots, K$
- Let N be the number of data points

40

Evaluation of Clustering: Purity Score

- For each cluster k ,
 - Count the number of data points from each class
 - Consider the number of data points of most common class

$$\max_m |N_k \cap C_m|$$

$|N_k \cap C_m|$ is the number of data points in k^{th} cluster belonging to class m


- Take the sum over all clusters, k
- Divide by the total number of data points (N)

$$\text{Purity Score} = \frac{1}{N} \sum_{k=1}^K \max_m |N_k \cap C_m|$$

41

Illustration of Computing Purity Score


- Number of data points, $N = 25$
- number of classes, $M = 3$
- number of clusters, $K = 3$



42

Illustration of Computing Purity Score


- Number of data points, $N = 25$
- number of classes, $M = 3$
- number of clusters, $K = 3$
- *Cluster 1:*
 - Number of examples of **Blue Class** are more, i.e. **5**



43

Illustration of Computing Purity Score


- Number of data points, $N = 25$
- number of classes, $M = 3$
- number of clusters, $K = 3$
- *Cluster 1:*
 - Number of examples of **Blue Class** are more, i.e. **5**
- *Cluster 2:*
 - Number of examples of **Red Class** are more, i.e. **5**



44

Illustration of Computing Purity Score


- Number of data points, $N = 25$
- number of classes, $M = 3$
- number of clusters, $K = 3$
- *Cluster 1:*
 - Number of examples of **Blue Class** are more, i.e. **5**
- *Cluster 2:*
 - Number of examples of **Red Class** are more, i.e. **5**
- *Cluster 3:*
 - Number of examples of **Green Class** are more, i.e. **5**



45

Illustration of Computing Purity Score

- Number of data points, $N = 25$
- number of classes, $M = 3$
- number of clusters, $K = 3$
- *Cluster 1:*
 - Number of examples of **Blue Class** are more, i.e. **5**
- *Cluster 2:*
 - Number of examples of **Red Class** are more, i.e. **5**
- *Cluster 3:*
 - Number of examples of **Green Class** are more, i.e. **5**
- **Purity score: $(5+5+5)/25 = 0.60$**



46

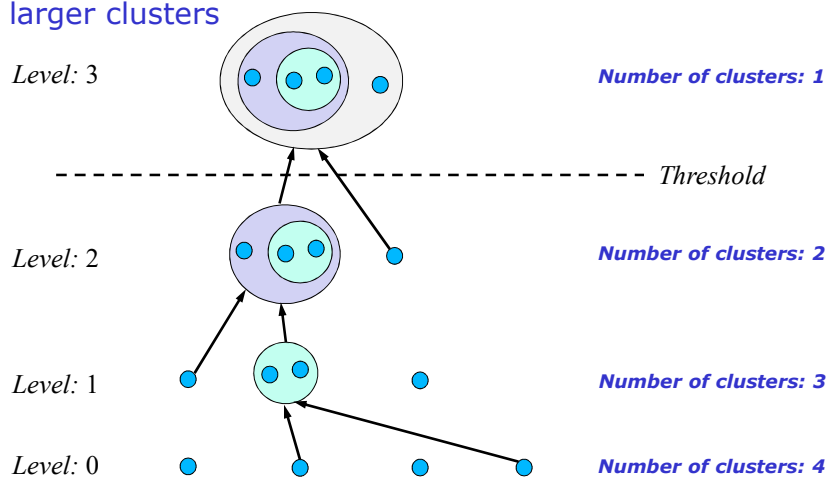
Hierarchical Clustering

Hierarchical Clustering Algorithms

- These methods create a **hierarchical decomposition** of the collection of examples
- Produce nested sequence of data partitions
- These sequence can be depicted using a tree structure
- Hierarchical clustering method works by grouping data points into a **tree of clusters**
- Hierarchical algorithms are either **agglomerative** or **divisive**
 - This classification of hierarchical clustering is depending on whether the hierarchical decomposition is formed in a
 - **Bottom-up (merging)** OR
 - **Top-down (splitting)** fashion
- Need not have to specify the number of clusters

Agglomerative Hierarchical Clustering

- Bottom-up approach
- This strategy starts by placing each example in its own cluster (atomic clusters or singleton clusters) and then merges these atomic clusters into larger and larger clusters



Agglomerative Hierarchical Clustering

- Bottom-up approach
- This strategy starts by placing each example in its own cluster (atomic clusters) and then merges these atomic clusters into larger and larger clusters
- Starts with N clusters where each example is a cluster
- At each successive step (level), the most similar pair of clusters are merged
 - The measure of closeness (intercluster similarity) is considered to decide which two clusters are merged
 - At each level, number of clusters is reduced by one
- The process continues till all the examples are in a single cluster or until certain termination conditions are satisfied
 - Termination condition could be
 - Number of clusters
 - Intercluster similarity between each pair of cluster is within a certain threshold

50

Agglomerative Hierarchical Clustering

- **Once two examples are placed in the same cluster at a level, they remain in same cluster at all subsequent levels**
- Example: AGglomerative NESTing (AGNES)
- Most hierarchical clustering methods belong to this category
- They differ only in their definition of intercluster similarity
- Intercluster similarity is to identifying two closest cluster for merging
- When there is one example in a cluster, two closest clusters are found by computing minimum Euclidian distance between two clusters
- However, there is no unique way to find the two closest clusters when there are more than one data points in each clusters

51

Agglomerative Hierarchical Clustering

- Different intercluster similarities to find similarity between the clusters having more than one examples:
 1. Minimum distance between any two examples from two clusters C_i and C_j

$$d_{\min}(C_i, C_j) = \min_{\mathbf{x} \in C_i, \mathbf{x}' \in C_j} \|\mathbf{x} - \mathbf{x}'\|$$

- Select a pair of clusters for merging whose minimum distance between any two examples is minimum of all the pair of clusters

2. Distance between the centers of two clusters C_i and C_j

$$d_{\text{mean}}(C_i, C_j) = \|\boldsymbol{\mu}_i - \boldsymbol{\mu}_j\|$$

- Where $\boldsymbol{\mu}_i$ is the center of C_i and $\boldsymbol{\mu}_j$ is the center of C_j
- Select a pair of clusters for merging whose distance between the centers is minimum of all the pair of clusters

52

Agglomerative Hierarchical Clustering

- Different **intercluster similarities** to find similarity between the clusters having **more than one examples**:
- 3. **Average distance** of all the points in one cluster (C_i) to all the points in another cluster (C_j)

$$d_{avg}(C_i, C_j) = \frac{1}{N_i N_j} \sum_{\mathbf{x} \in C_i} \sum_{\mathbf{x}' \in C_j} \|\mathbf{x} - \mathbf{x}'\|$$

- Where N_i and N_j are the number of examples in clusters C_i and C_j respectively
- Select a pair of clusters for merging whose **average distance** is **minimum than that of all the pair of clusters**

53


Agglomerative Hierarchical Clustering

- **Given**: Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^d$
- **Target**: Partition the data
- **Step1**: N clusters where each example is a cluster
- **Step2**: Compute **intercluster similarity** between each pair of clusters
- **Step3**: Choose a pair of clusters that are **most similar** (**minimum intercluster distance**) and merge them
- **Step4**: Repeat **Step2** and **Step3** until **all the examples are in a single cluster** or until **certain termination conditions are satisfied**

54

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 0

Number of clusters: 10

55

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 1

Number of clusters: 9

56

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 2

Number of clusters: 8

57

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 3

Number of clusters: 7

58

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 4

Number of clusters: 6

59

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 5

Number of clusters: 5

60

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 6

Number of clusters: 4

61

Illustration: Agglomerative Hierarchical Clustering

- **Intercluster similarity:**
 - **Single example in clusters:** Euclidian distance
 - **More than one examples in cluster:** Distance between the centres of two clusters




Level: 7

Number of clusters: 3

62

Divisive Hierarchical Clustering

- Top-down approach
- Starts with single cluster having all the examples
- It subdivides the cluster into smaller and smaller clusters in the successive step



63

Divisive Hierarchical Clustering

- Top-down approach
- Starts with single cluster having all the examples
- It subdivides the cluster into smaller and smaller clusters in the successive step
- At each successive step, a compactness measure is used to choose which cluster to split
 - Compactness measure: Average value of distance between the data points of a cluster
 - Compactness measure (CM_i) of a cluster C_i :

$$CM_i = \frac{1}{N_i^2} \sum_{\mathbf{x} \in C_i} \sum_{\mathbf{x}' \in C_i} \|\mathbf{x} - \mathbf{x}'\|$$

- Where N_i is the number of examples in cluster C_i
- Choose the cluster with larger value of compactness measure to split

64

Divisive Hierarchical Clustering

- To split a cluster, find a pair of examples having maximum Euclidian distance and split around these two examples (keeping them as centroids)
- At each level, number of clusters is increases by one
- The process continues until each example forms a cluster (atomic or singleton cluster) or until it satisfies certain termination condition
 - Termination condition could be
 - Number of clusters
 - Compactness measure of each cluster is within a certain threshold
- **Once two examples are placed in two different clusters at a level, they remain in different clusters at all subsequent levels**
- Example: DDivisive ANALysis (DIANA)


65

Divisive Hierarchical Clustering

- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^d$
- **Target:** Partition the data
- **Step1:** Single cluster having all the examples
- **Step2:** Find a pair of examples with in a cluster having maximum Euclidian distance
 - These examples act as centroid
- **Step3:** Split into two clusters by assigning each data point to one of these two examples using Euclidian distance
- **Step4:** Compute compactness measure for each cluster
- **Step5:** Choose the cluster with larger value of compactness measure to split
- **Step6:** Repeat Step2 to Step5 until each example forms a cluster (atomic or singleton cluster) or until it satisfies certain termination condition

66

Divisive Hierarchical Clustering




Level: 0

Number of clusters: 1

67

Divisive Hierarchical Clustering




Level: 0

Number of clusters: 1

68

Divisive Hierarchical Clustering




Level: 1

Number of clusters: 2

69

Divisive Hierarchical Clustering




Level: 1

Number of clusters: 2

70

Divisive Hierarchical Clustering



Level: 2


Number of clusters: 3

71

Density-Based Clustering

Density-Based Clustering

- These methods cluster collection of examples based on the notion of **density**
- These methods regard **clusters as dense regions of examples** in the data space that are separated by regions of low density (i.e. noise)
- They discover clusters with **arbitrary shape**



73


Density-Based Clustering

- These methods cluster collection of examples based on the notion of **density**
- These methods regard **clusters as dense regions of examples** in the data space that are separated by regions of low density (i.e. noise)
- They discover clusters with **arbitrary shape**
- **They automatically identifies the number of clusters**
- **General idea:** To continue growing the given cluster as long as density (number of examples) in the neighbourhood exceeds some threshold
- **Example: Density-based Spatial Clustering of Applications with Noise (DBSCAN)**
 - It grows the clusters according to a density-based connectivity analysis

74

Density-based Spatial Clustering of Applications with Noise (DBSCAN)


- DBSCAN is a density-based clustering included with noise
- It grows the regions with sufficiently high density (neighbors) into clusters with arbitrary shape
- It defines a cluster as a maximal set of **density-connected points**
- DBSCAN has 5 important components:
 1. **Epsilon (ϵ)**: It is a value of radius of boundary from every example



75

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 1. **Epsilon (ϵ)**: It is a value of radius of boundary from every example
 2. **MinPoints**: **Minimum number of examples** present inside the boundary with radius of ϵ from an example x
 - These examples with in a boundary are neighbors to x and called as **ϵ -neighborhood of an example, x**



76

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 - Epsilon (ϵ)**: It is a value of radius of boundary from every example
 - MinPoints**: Minimum number of examples present inside the boundary with radius of ϵ from an example x
 - Core point**: If there are atleast **MinPoints** number of examples are with in ϵ -radius from x , then x is called as **core point**



77

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 - Epsilon (ϵ)**: It is a value of radius of boundary from every example
 - MinPoints**: Minimum number of examples present inside the boundary with radius of ϵ from an example x
 - Core point**: If there are atleast **MinPoints** number of examples are with in ϵ -radius from x , then x is called as **core point**



78

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 - Epsilon (ϵ)**: It is a value of radius of boundary from every example
 - MinPts**: Minimum number of examples present inside the boundary with radius of ϵ from an example x
 - Core point**: If there are atleast **MinPoints** number of examples are with in ϵ -radius from x , then x is called as **core point**



79

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 - Core point**: If there are atleast **MinPoints** number of examples are with in ϵ -radius from x , then x is called as **core point**
 - Border point**:
 - The number of examples within ϵ -radius from x is **less than MinPoints** **AND** atleast one of the example in neighborhood is **core point**, then x is called as **border point**



80

Density-based Special Clustering of Applications with Noise (DBSCAN)


- DBSCAN has 5 important components:
 3. **Core point**: If there are atleast *MinPoints* number of examples are with in ϵ -radius from x , then x is called as **core point**
 4. **Border point**:
 - The number of examples within ϵ -radius from x is **less than *MinPoints*** **AND** atleast one of the example in neighborhood is **core point**, then x is called as **border point**



81

Density-based Special Clustering of Applications with Noise (DBSCAN)

- DBSCAN has 5 important components:
 5. **Noise point**:
 - The number of examples within ϵ -radius from x is **less than *MinPoints*** **AND** no example in neighborhood is **core point**
 - The noise point is similar to outlier



82

Clustering using DBSCAN


- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^d$
- First step is to identify **core points**, **border points** and **noise points**
 - Only **core points** and **border points** are considered inside the cluster
 - **Noise points** are not taken into the cluster
 - Thus DBSCAN is robust to outliers
- Next step is to find the **connected components** of core points
- **Connected component of core points:** Connecting the core points that are reachable from any point
 - All the connected (reachable) core points form a cluster

Clustering using DBSCAN

- The **connected component of core points** is obtained by understanding following **two** definitions.
- **Directly density-reachable:** A core point \mathbf{x}_i is directly density-reachable to a core point \mathbf{x}_j , if the core point \mathbf{x}_j is within ϵ -distance from core point \mathbf{x}_i
- **Density-reachable:** A core point \mathbf{x}_i is indirectly reachable to another core point \mathbf{x}_j through other core points, $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_k, \mathbf{x}_{k+1}, \dots, \mathbf{x}_K$ such that
 - \mathbf{x}_i is directly density-reachable to \mathbf{x}_1
 - \mathbf{x}_1 is directly density-reachable to \mathbf{x}_2
 -
 - \mathbf{x}_k is directly density-reachable to \mathbf{x}_{k+1}
 -
 - \mathbf{x}_K is directly density-reachable to \mathbf{x}_j


Clustering using DBSCAN

- Directly density-reachable:
- Density-reachable:




Clustering using DBSCAN

- Directly density-reachable:
- Density-reachable:




Clustering using DBSCAN

- All the core points with connected component forms a cluster




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point



Clustering using DBSCAN

- All the core points with connected component forms a cluster
- Assign the **border points** to nearby cluster which is at ϵ -radius from that border point




Clustering using DBSCAN

- **Training process:**
- **Given:** Training data, $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^d$
- Identify the core points, border points and noise points
- Find the connected components of core points
- Each connected component forms a cluster
- Assign each of the border points to a nearby cluster which is at ϵ -radius from that border point
- Noise points are not assigned to any clusters
- Training process, stores the core points as model

Clustering using DBSCAN


- **Test process:**
- For a test example, identify it as **core point** or **border point** or **noise point**
- If it is a **core point**, assign it to a cluster to which it is **directly density-reachable** or **density-reachable**
- If it is a **border point**, assign it to a **nearby cluster** which is at ϵ -radius from that border point
- If it is a **noise point**, do not assign to any cluster

Clustering using DBSCAN




Advantages of DBSCAN

- DBSCAN **does not** require to specify the number of clusters in the data a priori
- DBSCAN can find **arbitrarily shaped clusters**
- DBSCAN has a notion of noise, and is robust to **outliers**
- The parameters ϵ and *MinPoints* are experimentally set by the users



Limitation of DBSCAN

- DBSCAN is not suitable when the data is completely dense and there is no low dense area to separate



- The parameters ϵ and *MinPoints* should be chosen carefully

Text Books

1. J. Han and M. Kamber, *Data Mining: Concepts and Techniques*, Third Edition, Morgan Kaufmann Publishers, 2011.
2. S. Theodoridis and K. Koutroumbas, *Pattern Recognition*, Academic Press, 2009.